

Hands-on Lab 7 Removing Duplicates_v2

March 22, 2025

1 Removing Duplicates

Estimated time needed: **30** minutes

1.1 Introduction

In this lab, you will focus on data wrangling, an important step in preparing data for analysis. Data wrangling involves cleaning and organizing data to make it suitable for analysis. One key task in this process is removing duplicate entries, which are repeated entries that can distort analysis and lead to inaccurate conclusions.

1.2 Objectives

In this lab you will perform the following:

1. Identify duplicate rows in the dataset.
2. Use suitable techniques to remove duplicate rows and verify the removal.
3. Summarize how to handle missing values appropriately.
4. Use `ConvertedCompYearly` to normalize compensation data.

1.2.1 Install the Required Libraries

```
[1]: !pip install pandas
```

```
Collecting pandas
```

```
  Downloading
```

```
pandas-2.2.3-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (89 kB)
```

```
Collecting numpy>=1.26.0 (from pandas)
```

```
  Downloading
```

```
numpy-2.2.4-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (62 kB)
```

```
Requirement already satisfied: python-dateutil>=2.8.2 in
```

```
/opt/conda/lib/python3.12/site-packages (from pandas) (2.9.0.post0)
```

```
Requirement already satisfied: pytz>=2020.1 in /opt/conda/lib/python3.12/site-packages (from pandas) (2024.2)
```

```
Collecting tzdata>=2022.7 (from pandas)
```

```
  Downloading tzdata-2025.1-py2.py3-none-any.whl.metadata (1.4 kB)
```

```
Requirement already satisfied: six>=1.5 in /opt/conda/lib/python3.12/site-packages (from python-dateutil>=2.8.2->pandas) (1.17.0)
```

Downloading
pandas-2.2.3-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (12.7 MB)
12.7/12.7 MB
130.3 MB/s eta 0:00:00
Downloading
numpy-2.2.4-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (16.1 MB)
16.1/16.1 MB
149.3 MB/s eta 0:00:00
Downloading tzdata-2025.1-py2.py3-none-any.whl (346 kB)
Installing collected packages: tzdata, numpy, pandas
Successfully installed numpy-2.2.4 pandas-2.2.3 tzdata-2025.1

1.2.2 Step 1: Import Required Libraries

```
[2]: import pandas as pd
```

1.2.3 Step 2: Load the Dataset into a DataFrame

load the dataset using pd.read_csv()

```
[17]: # Define the URL of the dataset
import pandas as pd

file_path = "https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/
↳n01PQ9pSmiRX6520flujwQ/survey-data.csv"

try:

    df = pd.read_csv(file_path, encoding='utf-8')
    print("Data loaded successfully!")
    print(df.head())
except Exception as e:
    print("Error loading data:", e)
```

Data loaded successfully!

	ResponseId	MainBranch	Age \
0	1	I am a developer by profession	Under 18 years old
1	2	I am a developer by profession	35-44 years old
2	3	I am a developer by profession	45-54 years old
3	4	I am learning to code	18-24 years old
4	5	I am a developer by profession	18-24 years old

	Employment	RemoteWork	Check \
0	Employed, full-time	Remote	Apples
1	Employed, full-time	Remote	Apples
2	Employed, full-time	Remote	Apples
3	Student, full-time	NaN	Apples

4	Student, full-time	NaN	Apples						
---	--------------------	-----	--------	--	--	--	--	--	--

				CodingActivities \
0				Hobby
1	Hobby;Contribute to open-source projects;Other...			
2	Hobby;Contribute to open-source projects;Other...			
3				NaN
4				NaN

				EdLevel \
0	Primary/elementary school			
1	Bachelor's degree (B.A., B.S., B.Eng., etc.)			
2	Master's degree (M.A., M.S., M.Eng., MBA, etc.)			
3	Some college/university study without earning ...			
4	Secondary school (e.g. American high school, G...			

				LearnCode \
0	Books / Physical media			
1	Books / Physical media;Colleague;On the job tr...			
2	Books / Physical media;Colleague;On the job tr...			
3	Other online resources (e.g., videos, blogs, f...			
4	Other online resources (e.g., videos, blogs, f...			

				LearnCodeOnline ...	JobSatPoints_6 \
0				NaN ...	NaN
1	Technical documentation;Blogs;Books;Written Tu... ...				0.0
2	Technical documentation;Blogs;Books;Written Tu... ...				NaN
3	Stack Overflow;How-to videos;Interactive tutorial ...				NaN
4	Technical documentation;Blogs;Written Tutorial... ...				NaN

	JobSatPoints_7	JobSatPoints_8	JobSatPoints_9	JobSatPoints_10 \
0	NaN	NaN	NaN	NaN
1	0.0	0.0	0.0	0.0
2	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN

	JobSatPoints_11	SurveyLength	SurveyEase	ConvertedCompYearly	JobSat
0	NaN	NaN	NaN	NaN	NaN
1	0.0	NaN	NaN	NaN	NaN
2	NaN	Appropriate in length	Easy	NaN	NaN
3	NaN	Too long	Easy	NaN	NaN
4	NaN	Too short	Easy	NaN	NaN

[5 rows x 114 columns]

Note: If you are working on a local Jupyter environment, you can use the URL directly in the `pandas.read_csv()` function as shown below:

```
#df = pd.read_csv("https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/n01PQ9pSmiRX6520flu  
data.csv")
```

1.2.4 Step 3: Identifying Duplicate Rows

Task 1: Identify Duplicate Rows 1. Count the number of duplicate rows in the dataset. 2. Display the first few duplicate rows to understand their structure.

```
[8]: # Count the number of duplicate rows  
num_duplicates = df.duplicated().sum()  
print("Number of duplicate rows:", num_duplicates)  
  
# Display the first few duplicate rows  
duplicate_rows = df[df.duplicated()]  
print("First few duplicate rows:")  
print(duplicate_rows.head())
```

Number of duplicate rows: 0

First few duplicate rows:

Empty DataFrame

Columns: [ResponseId, MainBranch, Age, Employment, RemoteWork, Check, CodingActivities, EdLevel, LearnCode, LearnCodeOnline, TechDoc, YearsCode, YearsCodePro, DevType, OrgSize, PurchaseInfluence, BuyNewTool, BuildvsBuy, TechEndorse, Country, Currency, CompTotal, LanguageHaveWorkedWith, LanguageWantToWorkWith, LanguageAdmired, DatabaseHaveWorkedWith, DatabaseWantToWorkWith, DatabaseAdmired, PlatformHaveWorkedWith, PlatformWantToWorkWith, PlatformAdmired, WebframeHaveWorkedWith, WebframeWantToWorkWith, WebframeAdmired, EmbeddedHaveWorkedWith, EmbeddedWantToWorkWith, EmbeddedAdmired, MiscTechHaveWorkedWith, MiscTechWantToWorkWith, MiscTechAdmired, ToolsTechHaveWorkedWith, ToolsTechWantToWorkWith, ToolsTechAdmired, NEWCollabToolsHaveWorkedWith, NEWCollabToolsWantToWorkWith, NEWCollabToolsAdmired, OpSysPersonal use, OpSysProfessional use, OfficeStackAsyncHaveWorkedWith, OfficeStackAsyncWantToWorkWith, OfficeStackAsyncAdmired, OfficeStackSyncHaveWorkedWith, OfficeStackSyncWantToWorkWith, OfficeStackSyncAdmired, AISearchDevHaveWorkedWith, AISearchDevWantToWorkWith, AISearchDevAdmired, NEWSOSites, SOVisitFreq, SOAccount, SOPartFreq, SOHow, SOComm, AISelect, AISent, AIBen, AIAcc, AIComplex, AIToolCurrently Using, AIToolInterested in Using, AIToolNot interested in Using, AINextMuch more integrated, AINextNo change, AINextMore integrated, AINextLess integrated, AINextMuch less integrated, AIThreat, AIEthics, AICHallenges, TBranch, ICorPM, WorkExp, Knowledge_1, Knowledge_2, Knowledge_3, Knowledge_4, Knowledge_5, Knowledge_6, Knowledge_7, Knowledge_8, Knowledge_9, Frequency_1, Frequency_2, Frequency_3, TimeSearching, TimeAnswering, Frustration, ProfessionalTech, ProfessionalCloud, ProfessionalQuestion, ...]

Index: []

[0 rows x 114 columns]

1.2.5 Step 4: Removing Duplicate Rows

Task 2: Remove Duplicates 1. Remove duplicate rows from the dataset using the `drop_duplicates()` function. 2. Verify the removal by counting the number of duplicate rows after removal.

```
[9]: ## Write your code here
      # Remove duplicate rows
      df = df.drop_duplicates()

      # Verify the removal by counting the number of duplicate rows
      num_duplicates = df.duplicated().sum()
      print("Number of duplicate rows after removal:", num_duplicates)
```

Number of duplicate rows after removal: 0

1.2.6 Step 5: Handling Missing Values

Task 3: Identify and Handle Missing Values 1. Identify missing values for all columns in the dataset. 2. Choose a column with significant missing values (e.g., `EdLevel`) and impute with the most frequent value.

```
[10]: ## Write your code here
       # Identify missing values for all columns
       missing_values = df.isnull().sum()
       print("Missing values per column:")
       print(missing_values)

       # Impute missing values in the 'EdLevel' column with the most frequent value
       most_frequent = df['EdLevel'].mode()[0]
       df['EdLevel'].fillna(most_frequent, inplace=True)

       # Verify the imputation
       print("\nMissing values in 'EdLevel' after imputation:")
       print(df['EdLevel'].isnull().sum())
```

Missing values per column:

ResponseId	0
MainBranch	0
Age	0
Employment	0
RemoteWork	10631
...	
JobSatPoints_11	35992
SurveyLength	9255
SurveyEase	9199
ConvertedCompYearly	42002
JobSat	36311

Length: 114, dtype: int64

Missing values in 'EdLevel' after imputation:

0

/tmp/ipykernel_300/462093557.py:9: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['EdLevel'].fillna(most_frequent, inplace=True)
```

1.2.7 Step 6: Normalizing Compensation Data

Task 4: Normalize Compensation Data Using ConvertedCompYearly 1. Use the ConvertedCompYearly column for compensation analysis as the normalized annual compensation is already provided. 2. Check for missing values in ConvertedCompYearly and handle them if necessary.

```
[11]: ## Write your code here
# Check for missing values in ConvertedCompYearly
missing_comp = df['ConvertedCompYearly'].isnull().sum()
print("Missing values in 'ConvertedCompYearly':", missing_comp)

# Handle missing values by imputing with the median
median_comp = df['ConvertedCompYearly'].median()
df['ConvertedCompYearly'].fillna(median_comp, inplace=True)

# Verify the imputation
missing_comp_after = df['ConvertedCompYearly'].isnull().sum()
print("Missing values in 'ConvertedCompYearly' after imputation:",
      missing_comp_after)
```

Missing values in 'ConvertedCompYearly': 42002

Missing values in 'ConvertedCompYearly' after imputation: 0

/tmp/ipykernel_300/2462540287.py:8: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value)

instead, to perform the operation inplace on the original object.

```
df['ConvertedCompYearly'].fillna(median_comp, inplace=True)
```

1.2.8 Step 7: Summary and Next Steps

In this lab, you focused on identifying and removing duplicate rows.

- You handled missing values by imputing the most frequent value in a chosen column.
- You used ConvertedCompYearly for compensation normalization and handled missing values.
- For further analysis, consider exploring other columns or visualizing the cleaned dataset.

```
[14]: pip install matplotlib seaborn
```

```
Collecting matplotlib
  Downloading matplotlib-3.10.1-cp312-cp312-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (11 kB)
Collecting seaborn
  Downloading seaborn-0.13.2-py3-none-any.whl.metadata (5.4 kB)
Collecting contourpy>=1.0.1 (from matplotlib)
  Downloading contourpy-1.3.1-cp312-cp312-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (5.4 kB)
Collecting cycler>=0.10 (from matplotlib)
  Downloading cycler-0.12.1-py3-none-any.whl.metadata (3.8 kB)
Collecting fonttools>=4.22.0 (from matplotlib)
  Downloading fonttools-4.56.0-cp312-cp312-
manylinux_2_5_x86_64.manylinux1_x86_64.manylinux_2_17_x86_64.manylinux2014_x86_6
4.whl.metadata (101 kB)
Collecting kiwisolver>=1.3.1 (from matplotlib)
  Downloading kiwisolver-1.4.8-cp312-cp312-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (6.2 kB)
Requirement already satisfied: numpy>=1.23 in /opt/conda/lib/python3.12/site-
packages (from matplotlib) (2.2.4)
Requirement already satisfied: packaging>=20.0 in
/opt/conda/lib/python3.12/site-packages (from matplotlib) (24.2)
Collecting pillow>=8 (from matplotlib)
  Downloading pillow-11.1.0-cp312-cp312-manylinux_2_28_x86_64.whl.metadata (9.1
kB)
Collecting pyparsing>=2.3.1 (from matplotlib)
  Downloading pyparsing-3.2.1-py3-none-any.whl.metadata (5.0 kB)
Requirement already satisfied: python-dateutil>=2.7 in
/opt/conda/lib/python3.12/site-packages (from matplotlib) (2.9.0.post0)
Requirement already satisfied: pandas>=1.2 in /opt/conda/lib/python3.12/site-
packages (from seaborn) (2.2.3)
Requirement already satisfied: pytz>=2020.1 in /opt/conda/lib/python3.12/site-
packages (from pandas>=1.2->seaborn) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in /opt/conda/lib/python3.12/site-
```

```

packages (from pandas>=1.2->seaborn) (2025.1)
Requirement already satisfied: six>=1.5 in /opt/conda/lib/python3.12/site-
packages (from python-dateutil>=2.7->matplotlib) (1.17.0)
Downloading
matplotlib-3.10.1-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl
(8.6 MB)
8.6/8.6 MB
140.0 MB/s eta 0:00:00
Downloading seaborn-0.13.2-py3-none-any.whl (294 kB)
Downloading
contourpy-1.3.1-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (323
kB)
Downloading cycler-0.12.1-py3-none-any.whl (8.3 kB)
Downloading fonttools-4.56.0-cp312-cp312-
manylinux_2_5_x86_64.manylinux1_x86_64.manylinux_2_17_x86_64.manylinux2014_x86_6
4.whl (4.9 MB)
4.9/4.9 MB
114.8 MB/s eta 0:00:00
Downloading
kiwisolver-1.4.8-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (1.5
MB)
1.5/1.5 MB
93.8 MB/s eta 0:00:00
Downloading pillow-11.1.0-cp312-cp312-manylinux_2_28_x86_64.whl (4.5 MB)
4.5/4.5 MB
157.4 MB/s eta 0:00:00
Downloading pyparsing-3.2.1-py3-none-any.whl (107 kB)
Installing collected packages: pyparsing, pillow, kiwisolver, fonttools, cycler,
contourpy, matplotlib, seaborn
Successfully installed contourpy-1.3.1 cycler-0.12.1 fonttools-4.56.0
kiwisolver-1.4.8 matplotlib-3.10.1 pillow-11.1.0 pyparsing-3.2.1 seaborn-0.13.2
Note: you may need to restart the kernel to use updated packages.

```

```
[15]: import matplotlib.pyplot as plt
import seaborn as sns
```

```
[16]: ## Write your code here

print("Summary statistics:")
print(df.describe())

# Visualize the distribution of annual compensation
plt.figure(figsize=(10, 6))
sns.histplot(df['ConvertedCompYearly'], bins=50, kde=True)
plt.title("Distribution of Annual Compensation")
plt.xlabel("Annual Compensation (USD)")
plt.ylabel("Frequency")

```



```
plt.show()

# Visualize the relationship between education level and compensation
plt.figure(figsize=(12, 6))
sns.boxplot(data=df, x='EdLevel', y='ConvertedCompYearly')
plt.title("Compensation by Education Level")
plt.xlabel("Education Level")
plt.ylabel("Annual Compensation (USD)")
plt.xticks(rotation=45)
plt.show()
```

Summary statistics:

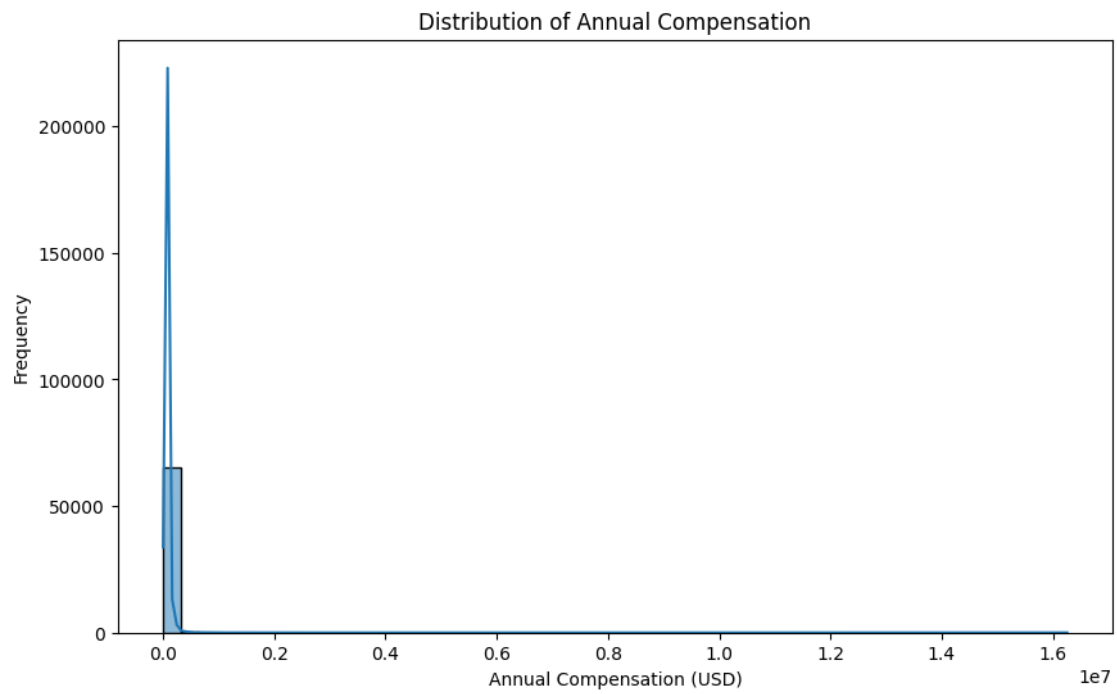
	ResponseId	CompTotal	WorkExp	JobSatPoints_1 \
count	65437.000000	3.374000e+04	29658.000000	29324.000000
mean	32719.000000	2.963841e+145	11.466957	18.581094
std	18890.179119	5.444117e+147	9.168709	25.966221
min	1.000000	0.000000e+00	0.000000	0.000000
25%	16360.000000	6.000000e+04	4.000000	0.000000
50%	32719.000000	1.100000e+05	9.000000	10.000000
75%	49078.000000	2.500000e+05	16.000000	22.000000
max	65437.000000	1.000000e+150	50.000000	100.000000


	JobSatPoints_4	JobSatPoints_5	JobSatPoints_6	JobSatPoints_7 \
count	29393.000000	29411.000000	29450.000000	29448.000000
mean	7.522140	10.060857	24.343232	22.96522
std	18.422661	21.833836	27.089360	27.01774
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	20.000000	15.000000
75%	5.000000	10.000000	30.000000	30.000000
max	100.000000	100.000000	100.000000	100.000000

	JobSatPoints_8	JobSatPoints_9	JobSatPoints_10	JobSatPoints_11 \
count	29456.000000	29456.000000	29450.000000	29445.000000
mean	20.278165	16.169432	10.955713	9.953948
std	26.108110	24.845032	22.906263	21.775652
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	10.000000	5.000000	0.000000	0.000000
75%	25.000000	20.000000	10.000000	10.000000
max	100.000000	100.000000	100.000000	100.000000

	ConvertedCompYearly	JobSat
count	6.543700e+04	29126.000000
mean	7.257636e+04	6.935041
std	1.122207e+05	2.088259
min	1.000000e+00	0.000000
25%	6.500000e+04	6.000000

50%	6.500000e+04	7.000000
75%	6.500000e+04	8.000000
max	1.625660e+07	10.000000





<!-- ## Change Log

| Date (YYYY-MM-DD) | Version | Changed By | Change Description |
|-------------------|---------|-------------------|--------------------|
| 2024-11-05 | 1.2 | Madhusudhan Moole | Updated lab |
| 2024-09-24 | 1.1 | Madhusudhan Moole | Updated lab |
| 2024-09-23 | 1.0 | Raghul Ramesh | Created lab |

!>

Copyright © IBM Corporation. All rights reserved.